

SQL Server 2008 Transact-SQL Recipes



Joseph Sack

SQL Server 2008 Transact-SQL Recipes

Copyright © 2008 by Joseph Sack

All rights reserved. No part of this work may be reproduced or transmitted in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval system, without the prior written permission of the copyright owner and the publisher.

ISBN-13 (pbk): 978-1-59059-980-8

ISBN-10 (pbk): 1-59059-980-2

ISBN-13 (electronic): 978-1-4302-0626-2

Printed and bound in the United States of America 9 8 7 6 5 4 3 2 1

Trademarked names may appear in this book. Rather than use a trademark symbol with every occurrence of a trademarked name, we use the names only in an editorial fashion and to the benefit of the trademark owner, with no intention of infringement of the trademark.

Lead Editor: Jonathan Gennick

Technical Reviewer: Evan Terry

Editorial Board: Clay Andres, Steve Anglin, Ewan Buckingham, Tony Campbell, Gary Cornell,

Jonathan Gennick, Matthew Moodie, Joseph Ottinger, Jeffrey Pepper, Frank Pohlmann,

Ben Renow-Clarke, Dominic Shakeshaft, Matt Wade, Tom Welsh

Project Manager: Susannah Davidson Pfalzer

Copy Editor: Ami Knox

Associate Production Director: Kari Brooks-Copony

Production Editor: Laura Cheu

Compositor: Dina Quan

Proofreader: Liz Welch

Indexer: Brenda Miller

Artist: April Milne

Cover Designer: Kurt Krames

Manufacturing Director: Tom Debolski

Distributed to the book trade worldwide by Springer-Verlag New York, Inc., 233 Spring Street, 6th Floor, New York, NY 10013. Phone 1-800-SPRINGER, fax 201-348-4505, e-mail orders-ny@springer-sbm.com, or visit <http://www.springeronline.com>.

For information on translations, please contact Apress directly at 2855 Telegraph Avenue, Suite 600, Berkeley, CA 94705. Phone 510-549-5930, fax 510-549-5939, e-mail info@apress.com, or visit <http://www.apress.com>.

Apress and friends of ED books may be purchased in bulk for academic, corporate, or promotional use. eBook versions and licenses are also available for most titles. For more information, reference our Special Bulk Sales—eBook Licensing web page at <http://www.apress.com/info/bulksales>.

The information in this book is distributed on an “as is” basis, without warranty. Although every precaution has been taken in the preparation of this work, neither the author(s) nor Apress shall have any liability to any person or entity with respect to any loss or damage caused or alleged to be caused directly or indirectly by the information contained in this work.

Contents at a Glance

About the Author	xxv
About the Technical Reviewer	xxvii
Acknowledgments	xxix
Introduction	xxxix

CHAPTER 1	SELECT	1
CHAPTER 2	Perform, Capture, and Track Data Modifications	63
CHAPTER 3	Transactions, Locking, Blocking, and Deadlocking	115
CHAPTER 4	Tables	143
CHAPTER 5	Indexes	197
CHAPTER 6	Full-Text Search	217
CHAPTER 7	Views	239
CHAPTER 8	SQL Server Functions	257
CHAPTER 9	Conditional Processing, Control-of-Flow, and Cursors	307
CHAPTER 10	Stored Procedures	325
CHAPTER 11	User-Defined Functions and Types	343
CHAPTER 12	Triggers	373
CHAPTER 13	CLR Integration	401
CHAPTER 14	XML, Hierarchies, and Spatial Data	419
CHAPTER 15	Hints	449
CHAPTER 16	Error Handling	459
CHAPTER 17	Principals	475
CHAPTER 18	Securables, Permissions, and Auditing	501
CHAPTER 19	Encryption	547
CHAPTER 20	Service Broker	579

■ CHAPTER 21	Configuring and Viewing SQL Server Options	615
■ CHAPTER 22	Creating and Configuring Databases	621
■ CHAPTER 23	Database Integrity and Optimization	669
■ CHAPTER 24	Maintaining Database Objects and Object Dependencies	687
■ CHAPTER 25	Database Mirroring	697
■ CHAPTER 26	Database Snapshots	717
■ CHAPTER 27	Linked Servers and Distributed Queries	723
■ CHAPTER 28	Query Performance Tuning	739
■ CHAPTER 29	Backup and Recovery	789
■ INDEX		823

Contents

About the Author	xxv
About the Technical Reviewer	xxvii
Acknowledgments	xxix
Introduction	xxxi

CHAPTER 1	SELECT	1
	The Basic SELECT Statement	1
	Selecting Specific Columns from a Table	2
	Selecting Every Column for Every Row	3
	Selective Querying Using a Basic WHERE Clause	3
	Using the WHERE Clause to Specify Rows Returned in the Result Set	4
	Combining Search Conditions	4
	Negating a Search Condition	6
	Keeping Your WHERE Clause Unambiguous	6
	Using Operators and Expressions	7
	Using BETWEEN for Date Range Searches	9
	Using Comparisons	9
	Checking for NULL Values	10
	Returning Rows Based on a List of Values	11
	Using Wildcards with LIKE	11
	Declaring and Assigning Values to Variables	12
	Grouping Data	14
	Using the GROUP BY Clause	14
	Using GROUP BY ALL	15
	Selectively Querying Grouped Data Using HAVING	16
	Ordering Results	17
	Using the ORDER BY Clause	17
	Using the TOP Keyword with Ordered Results	19
	SELECT Clause Techniques	21
	Using DISTINCT to Remove Duplicate Values	21
	Using DISTINCT in Aggregate Functions	22
	Using Column Aliases	22
	Using SELECT to Create a Script	23
	Performing String Concatenation	24
	Creating a Comma-Delimited List Using SELECT	25
	Using the INTO Clause	26

Subqueries	27
Using Subqueries to Check for Matches	27
Querying from More Than One Data Source	28
Using INNER Joins	29
Using OUTER Joins	30
Using CROSS Joins	31
Referencing a Single Table Multiple Times in the Same Query	32
Using Derived Tables	33
Combining Result Sets with UNION	33
Using APPLY to Invoke a Table-Valued Function for Each Row	35
Using CROSS APPLY	35
Using OUTER APPLY	37
Advanced Techniques for Data Sources	38
Using the TABLESAMPLE to Return Random Rows	38
Using PIVOT to Convert Single Column Values into Multiple Columns and Aggregate Data	39
Normalizing Data with UNPIVOT	42
Returning Distinct or Matching Rows Using EXCEPT and INTERSECT	44
Summarizing Data	46
Summarizing Data Using CUBE	46
Summarizing Data Using ROLLUP	48
Creating Custom Summaries Using Grouping Sets	49
Revealing Rows Generated by GROUPING	51
Advanced Group-Level Identification with GROUPING_ID	53
Common Table Expressions	56
Using a Non-Recursive Common Table Expression	56
Using a Recursive Common Table Expression	59

■ CHAPTER 2 **Perform, Capture, and Track Data Modifications**..... 63

INSERT	63
Inserting a Row into a Table	64
Inserting a Row Using Default Values	65
Explicitly Inserting a Value into an IDENTITY Column	66
Inserting a Row into a Table with a uniqueidentifier Column	67
Inserting Rows Using an INSERT...SELECT Statement	68
Inserting Data from a Stored Procedure Call	70
Inserting Multiple Rows with VALUES	71
Using VALUES As a Table Source	72
UPDATE	73
Updating a Single Row	74
Updating Rows Based on a FROM and WHERE Clause	75
Updating Large Value Data Type Columns	76
Inserting or Updating an Image File Using OPENROWSET and BULK	78

Storing Unstructured Data on the File System While Maintaining SQL Server Transactional Control	80
Assigning and Modifying Database Values “in Place”	84
DELETE	86
Deleting Rows	86
Truncating a Table	88
Advanced Data Modification Techniques	89
Chunking Data Modifications with TOP	89
Executing INSERTs, UPDATEs, and DELETEs in a Single Statement	90
Capturing and Tracking Data Modification Changes	93
Returning Rows Affected by a Data Modification Statement	93
Asynchronously Capturing Table Data Modifications	96
Querying All Changes from CDC Tables	99
Querying Net Changes from CDC Tables	103
Translating the CDC Update Mask	104
Working with LSN Boundaries	105
Disabling Change Data Capture from Tables and the Database	107
Tracking Net Data Changes with Minimal Disk Overhead	107
 CHAPTER 3 Transactions, Locking, Blocking, and Deadlocking	 115
Transaction Control	115
Using Explicit Transactions	117
Displaying the Oldest Active Transaction with DBCC OPENTRAN	119
Querying Transaction Information by Session	120
Locking	122
Viewing Lock Activity	124
Controlling a Table’s Lock Escalation Behavior	126
Transaction, Locking, and Concurrency	128
Configuring a Session’s Transaction Locking Behavior	129
Blocking	134
Identifying and Resolving Blocking Issues	134
Configuring How Long a Statement Will Wait for a Lock to Be Released	136
Deadlocking	137
Identifying Deadlocks with a Trace Flag	138
Setting Deadlock Priority	141
 CHAPTER 4 Tables	 143
Table Basics	143
Creating a Table	147
Adding a Column to an Existing Table	147
Changing an Existing Column Definition	148

Creating a Computed Column	149
Reducing Storage for Null Columns	150
Dropping a Table Column	153
Reporting Table Information	154
Dropping a Table	154
Collation Basics	155
Viewing Collation Metadata	155
Designating a Column's Collation	156
Keys	157
Creating a Table with a Primary Key	158
Adding a Primary Key Constraint to an Existing Table	159
Creating a Table with a Foreign Key Reference	160
Adding a Foreign Key to an Existing Table	161
Creating Recursive Foreign Key References	162
Allowing Cascading Changes in Foreign Keys	163
Surrogate Keys	165
Using the IDENTITY Property During Table Creation	165
Using DBCC CHECKIDENT to View and Correct IDENTITY Seed Values	166
Using the ROWGUIDCOL Property	168
Constraints	168
Creating a Unique Constraint	169
Adding a UNIQUE Constraint to an Existing Table	170
Using CHECK Constraints	171
Adding a CHECK Constraint to an Existing Table	172
Disabling and Enabling a Constraint	173
Using a DEFAULT Constraint During Table Creation	174
Adding a DEFAULT Constraint to an Existing Table	175
Dropping a Constraint from a Table	176
Temporary Tables and Table Variables	176
Using a Temporary Table for Multiple Lookups Within a Batch	177
Creating a Table Variable to Hold a Temporary Result Set	178
Manageability for Very Large Tables	180
Implementing Table Partitioning	181
Determining the Location of Data in a Partition	184
Adding a New Partition	186
Removing a Partition	188
Moving a Partition to a Different Table	189
Removing Partition Functions and Schemes	190
Easing VLDB Manageability with Filegroups	191
Reducing Disk Space Usage with Data Compression	192

CHAPTER 5	Indexes	197
	Index Overview	197
	Creating a Table Index	199
	Enforcing Uniqueness on Non-Key Columns	201
	Creating an Index on Multiple Columns	202
	Defining Index Column Sort Direction	203
	Viewing Index Meta Data	203
	Disabling an Index	205
	Dropping Indexes	206
	Changing an Existing Index with DROP_EXISTING	206
	Controlling Index Build Performance and Concurrency	207
	Intermediate Index Creation in Tempdb	207
	Controlling Parallel Plan Execution for Index Creation	208
	Allowing User Table Access During Index Creation	208
	Index Options	209
	Using an Index INCLUDE	209
	Using PAD_INDEX and FILLFACTOR	210
	Disabling Page and/or Row Index Locking	211
	Managing Very Large Indexes	212
	Creating an Index on a Filegroup	212
	Implementing Index Partitioning	213
	Indexing a Subset of Rows	214
	Reducing Index Size	215
CHAPTER 6	Full-Text Search	217
	Full-Text Indexes and Catalogs	217
	Creating a Full-Text Catalog	217
	Creating a Full-Text Index	219
	Modifying a Full-Text Catalog	221
	Modifying a Full-Text Index	222
	Retrieving Full-Text Catalog and Index Metadata	225
	Discarding Common Strings from a Full-Text Index	226
	Dropping a Full-Text Index	229
	Dropping a Full-Text Catalog	230
	Basic Searching	230
	Using FREETEXT to Search Full-Text Indexed Columns	231
	Using CONTAINS for Word Searching	232
	Advanced Searching	232
	Using CONTAINS to Search with Wildcards	233
	Using CONTAINS to Search for Inflectional Matches	233
	Using CONTAINS for Searching Results by Term Proximity	234

Ranked Searching	235
Returning Ranked Search Results by Meaning.	235
Returning Ranked Search Results by Weighted Value.	236

■ CHAPTER 7 **Views** 239

Regular Views.	240
Creating a Basic View	240
Querying the View Definition	242
Displaying Views and Their Structures	243
Refreshing a View's Definition	244
Modifying a View	245
Dropping a View.	245
Modifying Data Through a View	246
View Encryption	247
Encrypting a View	247
Indexed Views	248
Creating an Indexed View	248
Forcing the Optimizer to Use an Index for an Indexed View	251
Partitioned Views	251
Creating a Distributed-Partitioned View	252

■ CHAPTER 8 **SQL Server Functions** 257

Aggregate Functions	257
Returning the Average of Values	258
Returning Row Counts.	259
Finding the Lowest and Highest Values from an Expression.	259
Returning the Sum of Values	260
Using Statistical Aggregate Functions.	260
Mathematical Functions	261
Performing Mathematical Operations	262
String Functions.	263
Converting a Character Value to ASCII and Back to Character	264
Returning Integer and Character Unicode Values.	265
Finding the Start Position of a String Within Another String	266
Finding the Start Position of a String Within Another String Using Wildcards	266
Determining the Similarity of Strings.	267
Taking the Leftmost or Rightmost Part of a String.	268
Determining the Number of Characters or Bytes in a String.	269
Replacing a Part of a String	269
Stuffing a String into a String	270
Changing Between Lower- and Uppercase.	270

Removing Leading and Trailing Blanks	271
Repeating an Expression N Number of Times	272
Repeating a Blank Space N Number of Times	272
Outputting an Expression in Reverse Order	273
Returning a Chunk of an Expression	273
Working with NULLs	274
Replacing a NULL Value with an Alternative Value.	274
Performing Flexible Searches Using ISNULL.	275
Returning the First Non-NULL Value in a List of Expressions	276
Returning a NULL Value When Two Expressions Are Equal: Otherwise Returning the First Expression	277
Date Functions	277
Returning the Current Date and Time	278
Converting Between Time Zones	279
Incrementing or Decrementing a Date's Value.	280
Finding the Difference Between Two Dates	281
Displaying the String Value for Part of a Date.	282
Displaying the Integer Representation for Parts of a Date.	282
Displaying the Integer Value for Part of a Date Using YEAR, MONTH, and DAY.	283
Type Conversion.	284
Converting Between Data Types	284
Converting Dates to Their Textual Representation.	285
Representing Binary Data in String Literals	286
Evaluating the Data Type Returned by an Expression	287
Ranking Functions	288
Generating an Incrementing Row Number	289
Returning Rows by Rank.	290
Returning Rows by Rank Without Gaps.	292
Using NTILE	292
Probing Server, Database, and Connection-Level Settings Using System Functions.	293
Determining the First Day of the Week	293
Viewing the Language Used in the Current Session	294
Viewing and Setting Current Connection Lock Timeout Settings	295
Displaying the Nesting Level for the Current Stored Procedure Context	295
Returning the Current SQL Server Instance Name and SQL Server Version	296
Returning the Current Connection's Session ID (SPID)	296
Returning the Number of Open Transactions	297
Retrieving the Number of Rows Affected by the Previous Statement	297
Retrieving System Statistics.	298

Displaying Database and SQL Server Settings	299
Returning the Current Database ID and Name	300
Returning a Database Object Name and ID	301
Returning the Application and Host for the Current User Session	301
Reporting Current User and Login Context	302
Viewing User Connection Options	303
IDENTITY and uniqueidentifier Functions	303
Returning the Last Identity Value	304
Returning an Identity Column's Seed and Incrementing Value	305
Creating a New uniqueidentifier Value	305

■ CHAPTER 9 **Conditional Processing, Control-of-Flow, and Cursors** 307

Conditional Processing	307
Using CASE to Evaluate a Single Input Expression	308
Using CASE to Evaluate Boolean Expressions	309
Using IF...ELSE	310
Control-of-Flow	312
Using RETURN	313
Using WHILE	314
Using GOTO	316
Using WAITFOR	318
Cursors	319
Creating and Using Transact-SQL Cursors	321

■ CHAPTER 10 **Stored Procedures** 325

Stored Procedure Basics	325
Creating a Basic Stored Procedure	326
Creating a Parameterized Stored Procedure	328
Using OUTPUT Parameters	330
Modifying a Stored Procedure	332
Dropping Stored Procedures	332
Executing Stored Procedures Automatically at SQL Server Startup	333
Reporting Stored Procedure Metadata	334
Documenting Stored Procedures	335
Stored Procedure Security	335
Encrypting a Stored Procedure	336
Using EXECUTE AS to Specify the Procedure's Security Context	337
Recompilation and Caching	340
RECOMPILE(ing) a Stored Procedure Each Time It Is Executed	341
Flushing the Procedure Cache	342

CHAPTER 11	User-Defined Functions and Types	343
	UDF Basics	343
	Creating Scalar User-Defined Functions	344
	Creating Inline User-Defined Functions	349
	Creating Multi-Statement User-Defined Functions	351
	Modifying User-Defined Functions	354
	Viewing UDF Metadata	356
	Dropping User-Defined Functions	356
	Benefitting from UDFs	357
	Maintaining Reusable Code	357
	Cross-Referencing Natural Key Values	359
	Replacing Views with Multi-Statement UDFs	362
	UDT Basics	365
	Creating and Using User-Defined Types	365
	Identifying Columns and Parameters with Dependencies	
	on User-Defined Types	367
	Dropping User-Defined Types	368
	Passing Table-Valued Parameters	369
CHAPTER 12	Triggers	373
	DML Triggers	374
	Creating an AFTER DML Trigger	375
	Creating an INSTEAD OF DML Trigger	378
	Handling Transactions Within DML Triggers	381
	Controlling DML Triggers Based on Modified Columns	384
	Viewing DML Trigger Metadata	385
	DDL Triggers	386
	Creating a DDL Trigger That Audits Database-Level Events	387
	Creating a DDL Trigger That Audits Server-Level Events	389
	Using a Logon Trigger	390
	Viewing DDL Trigger Metadata	392
	Managing Triggers	393
	Modifying a Trigger	393
	Enabling and Disabling Table Triggers	394
	Limiting Trigger Nesting	395
	Controlling Trigger Recursion	396
	Setting Trigger Firing Order	397
	Dropping a Trigger	399

CHAPTER 13	CLR Integration	401
	CLR Overview	402
	When (and When Not) to Use Assemblies	402
	CLR Objects Overview	404
	Creating CLR Database Objects	404
	Enabling CLR Support in SQL Server	405
	Writing an Assembly for a CLR Stored Procedure	405
	Compiling an Assembly into a DLL File	408
	Loading the Assembly into SQL Server	409
	Creating the CLR Stored Procedure	410
	Creating a CLR Scalar User-Defined Function	412
	Creating a CLR Trigger	415
	Administering Assemblies	417
	Viewing Assembly Metadata	417
	Modifying an Assembly's Permissions	417
	Removing an Assembly from the Database	418
CHAPTER 14	XML, Hierarchies, and Spatial Data	419
	Working with Native XML	419
	Creating XML Data Type Columns	419
	Inserting XML Data into a Column	421
	Validating XML Data Using Schemas	422
	Retrieving XML Data	424
	Modifying XML Data	427
	Indexing XML Data	428
	Converting Between XML Documents and Relational Data	430
	Formatting Relational Data As XML	430
	Converting XML to a Relational Form	433
	Working with Native Hierarchical Data	435
	Storing Hierarchical Data	435
	Returning a Specific Ancestor	438
	Returning Child Nodes	439
	Returning a Node's Depth	440
	Returning the Root Node	440
	Determining Whether a Node Is a Child of the Current Node	441
	Changing Node Locations	441
	Native Spatial Data	442
	Storing Spatial Data	442
	Querying Spatial Data	445

CHAPTER 15	Hints	449
	Using Join Hints	449
	Forcing a HASH Join	450
	Using Query Hints	451
	Forcing a Statement Recompile	452
	Using Table Hints	454
	Executing a Query Without Locking	456
	Forcing a SEEK over a SCAN	456
CHAPTER 16	Error Handling	459
	System-Defined and User-Defined Error Messages	459
	Viewing System Error Information	459
	Creating a User-Defined Error Message	460
	Dropping a User-Defined Error Message	462
	Manually Raising an Error	462
	Invoking an Error Message	463
	Trapping and Handling Application Errors	465
	Old-Style Error Handling	466
	Error Handling with TRY...CATCH	468
	Applying Error Handling Without Recoding a Stored Procedure	470
	Nesting Error Handling	471
CHAPTER 17	Principals	475
	Windows Principals	475
	Creating a Windows Login	476
	Viewing Windows Logins	477
	Altering a Windows Login	478
	Dropping a Windows Login	479
	Denying SQL Server Access to a Windows User or Group	480
	SQL Server Principals	480
	Creating a SQL Server Login	482
	Viewing SQL Server Logins	482
	Altering a SQL Server Login	483
	Managing a Login's Password	484
	Dropping a SQL Login	485
	Managing Server Role Members	485
	Reporting Fixed Server Role Information	486
	Database Principals	488
	Creating Database Users	489
	Reporting Database User Information	490
	Modifying a Database User	490

Removing a Database User from the Database	491
Fixing Orphaned Database Users	491
Reporting Fixed Database Roles Information	493
Managing Fixed Database Role Membership	494
Managing User-Defined Database Roles	495
Managing Application Roles	497
 ■ CHAPTER 18 Securables, Permissions, and Auditing	501
Permissions Overview	502
Reporting SQL Server Assignable Permissions	503
Server-Scoped Securables and Permissions	505
Managing Server Permissions	507
Querying Server-Level Permissions	508
Database-Scoped Securables and Permissions	509
Managing Database Permissions	510
Querying Database Permissions	511
Schema-Scoped Securables and Permissions	514
Managing Schemas	516
Managing Schema Permissions	517
Object Permissions	519
Managing Object Permissions	521
Managing Permissions Across Securable Scopes	522
Determining a Current Connection's Permissions to a Securable	522
Reporting the Permissions for a Principal by Securable Scope	523
Changing Securable Ownership	527
Allowing SQL Logins to Access Non-SQL Server Resources	528
Auditing SQL Instance and Database-Level Activity of Principals	
Against Securables	529
Defining Audit Data Sources	530
Capturing SQL Instance-Scoped Events	533
Capturing Database-Scoped Events	535
Querying Captured Audit Data	539
Managing, Modifying, and Removing Audit Objects	543
 ■ CHAPTER 19 Encryption	547
Encryption by Passphrase	547
Using a Function to Encrypt by Passphrase	548
Master Keys	550
Backing Up and Restoring a Service Master Key	550
Creating, Regenerating, and Dropping a Database Master Key	551

Backing Up and Restoring a Database Master Key	553
Removing Service Master Key Encryption from the Database Master Key	554
Asymmetric Key Encryption	555
Creating an Asymmetric Key	555
Viewing Asymmetric Keys in the Current Database	556
Modifying the Asymmetric Key's Private Key Password	557
Encrypting and Decrypting Data Using an Asymmetric Key	557
Dropping an Asymmetric Key	560
Symmetric Key Encryption	560
Creating a Symmetric Key	560
Viewing Symmetric Keys in the Current Database	562
Changing How a Symmetric Key Is Encrypted	562
Using Symmetric Key Encryption and Decryption	563
Dropping a Symmetric Key	567
Certificate Encryption	567
Creating a Database Certificate	567
Viewing Certificates in the Database	568
Backing Up and Restoring a Certificate	569
Managing a Certificate's Private Key	570
Using Certificate Encryption and Decryption	571
Automatically Opening and Decrypting via a Symmetric Key	573
Transparent Data Encryption	575
Enabling Transparent Data Encryption	575
Managing and Removing TDE	576

■ CHAPTER 20 **Service Broker** 579

Example Scenario: Online Bookstore	580
Creating a Basic Service Broker Application	580
Enabling Databases for Service Broker Activity	581
Creating the Database Master Key for Encryption	582
Managing Message Types	582
Creating Contracts	584
Creating Queues	585
Creating Services	587
Initiating a Dialog	589
Querying the Queue for Incoming Messages	591
Receiving and Responding to a Message	591
Ending a Conversation	594
Prioritizing Service Broker Conversations	596

Creating a Stored Procedure to Process Messages	598
Creating the Bookstore Stored Procedure	598
Remote-Server Service Broker Implementations	601
Enabling Transport Security	603
Enabling Dialog Security	606
Creating Routes and Remote Service Bindings	608
Event Notifications	612
Capturing Login Commands	612
 ■ CHAPTER 21	
Configuring and Viewing SQL Server Options	615
Viewing SQL Server Configurations	615
Changing SQL Server Configurations	617
 ■ CHAPTER 22	
Creating and Configuring Databases	621
Creating, Altering, and Dropping Databases	621
Creating a Database with a Default Configuration	622
Viewing Database Information	622
Creating a Database Using File Options	624
Creating a Database with a User-Defined Filegroup	627
Setting Database User Access	628
Renaming a Database	631
Dropping a Database	632
Detaching a Database	632
Attaching a Database	634
Configuring Database Options	635
Viewing Database Options	636
Configuring ANSI SQL Options	636
Configuring Automatic Options	638
Creating or Modifying a Database to Allow External Access	640
Creating or Changing a Database to Use a Non-Server Default Collation	641
Configuring Cursor Options	642
Enabling Date Correlation Optimization	643
Modifying Database Parameterization Behavior	644
Enabling Read Consistency for a Transaction	647
Configuring Database Recovery Models	649
Configuring Page Verification	650
Controlling Database Access and Ownership	651
Changing a Database State to Online, Offline, or Emergency	652
Changing a Database Owner	653

Managing Database Files and Filegroups	654
Adding a Data File or Log File to an Existing Database	654
Removing a Data or Log File from a Database	656
Relocating a Data or Transaction Log File	657
Changing a File's Logical Name	658
Increasing a Database's File Size and Modifying Its Growth Options	659
Adding a Filegroup to an Existing Database	660
Setting the Default Filegroup	660
Removing a Filegroup	661
Making a Database or Filegroup Read-Only	662
Viewing and Managing Database Space Usage	663
Viewing Database Space Usage	663
Shrinking the Database or a Database File	665

■ CHAPTER 23 **Database Integrity and Optimization** 669

Database Integrity Checking	669
Checking Consistency of the Disk Space Allocation Structures with DBCC CHECKALLOC	670
Checking Allocation and Structural Integrity with DBCC CHECKDB	672
Tables and Constraints	674
Checking Allocation and Structural Integrity of All Tables in a Filegroup Using DBCC CHECKFILEGROUP	675
Checking Data Integrity for Tables and Indexed Views Using DBCC CHECKTABLE	676
Checking Table Integrity with DBCC CHECKCONSTRAINTS	679
Checking System Table Consistency with DBCC CHECKCATALOG	681
Index Maintenance	682
Rebuilding Indexes	682
Defragmenting Indexes	685
Rebuilding a Heap	686

■ CHAPTER 24 **Maintaining Database Objects and Object Dependencies** 687

Database Object Maintenance	687
Changing the Name of a User-Created Database Object	687
Changing an Object's Schema	689
Object Dependencies	690
Identifying Object Dependencies	690
Identifying Referencing and Referenced Entities	692
Viewing an Object's Definition	694

CHAPTER 25	Database Mirroring	697
	Database Mirroring in Context	697
	Database Mirroring Architecture	698
	Setting Up Database Mirroring	700
	Creating Mirroring Endpoints	700
	Backing Up and Restoring Principal Databases	705
	Creating a Database Mirroring Session	707
	Setup Summary	709
	Operating Database Mirroring	710
	Changing Operating Modes	711
	Performing Failovers	712
	Pausing or Resuming a Mirroring Session	713
	Stopping Mirroring Sessions and Removing Endpoints	714
	Monitoring and Configuring Options	714
	Monitoring Mirror Status	714
	Configuring the Connection Timeout Period	715
CHAPTER 26	Database Snapshots	717
	Snapshot Basics	717
	Creating and Querying Database Snapshots	718
	Removing a Database Snapshot	719
	Recovering Data with a Database Snapshot	720
CHAPTER 27	Linked Servers and Distributed Queries	723
	Linked Server Basics	723
	Creating a Linked Server to Another SQL Server Instance	724
	Configuring Linked Server Properties	725
	Viewing Linked Server Information	727
	Dropping a Linked Server	727
	Linked Server Logins	728
	Adding a Linked Server Login Mapping	728
	Viewing Linked Logins	729
	Dropping a Linked Server Login Mapping	730
	Executing Distributed Queries	730
	Executing Distributed Queries Against a Linked Server	730
	Creating and Using an Alias to Reference Four-Part Linked Server Names	732
	Executing Distributed Queries Using OPENQUERY	733
	Executing Ad Hoc Queries Using OPENROWSET	733
	Reading Data from a File Using OPENROWSET BULK Options	735

CHAPTER 28	Query Performance Tuning	739
	Query Performance Tips	740
	Capturing and Evaluating Query Performance	742
	Capturing Executing Queries Using sys.dm_exec_requests	742
	Viewing Estimated Query Execution Plans Using	
	Transact-SQL Commands	743
	Viewing Execution Runtime Information	746
	Viewing Performance Statistics for Cached Query Plans	748
	Viewing Aggregated Performance Statistics Based on Query	
	or Plan Patterns	750
	Identifying the Top Bottleneck	752
	Identifying I/O Contention by Database and File	753
	Index Tuning	754
	Displaying Index Fragmentation	756
	Displaying Index Usage	759
	Statistics	760
	Manually Creating Statistics	760
	Creating Statistics on a Subset of Rows	761
	Updating Statistics	762
	Generating and Updating Statistics Across All Tables	763
	Viewing Statistics Details	765
	Removing Statistics	766
	Miscellaneous Techniques	766
	Using an Alternative to Dynamic SQL	767
	Forcing SQL Server to Use a Query Plan	769
	Applying Hints Without Modifying Application SQL	771
	Creating Plan Guides from Cache	775
	Checking the Validity of a Plan Guide	777
	Parameterizing a Non-parameterized Query Using Plan Guides	778
	Limiting Competing Query Resource Consumption	781
CHAPTER 29	Backup and Recovery	789
	Creating a Backup and Recovery Plan	789
	Making Backups	791
	Performing a Basic Full Backup	793
	Compressing Your Backups	794
	Naming and Describing Your Backups and Media	796
	Configuring Backup Retention	797
	Striping Backup Sets	799
	Using a Named Backup Device	799
	Mirroring Backup Sets	801

- Performing a Transaction Log Backup 803
- Create Backups Without Breaking the Backup Sequence 804
- Performing a Differential Backup 805
- Backing Up Individual Files or Filegroups 805
- Performing a Partial Backup 807
- Viewing Backup Metadata 808
- Restoring a Database 810
 - Restoring a Database from a Full Backup 810
 - Restoring a Database from a Transaction Log Backup 812
 - Restoring a Database from a Differential Backup 815
 - Restoring a File or Filegroup 816
 - Performing a Piecemeal (PARTIAL) Restore 818
 - Restoring a Page 819
 - Identifying Databases with Multiple Recovery Paths 820

■ INDEX 823

About the Author



JOSEPH SACK is a dedicated support engineer in the Microsoft Premier Field Engineering organization and has worked with SQL Server since 1997. He is the author of *SQL Server 2005 T-SQL Recipes* (Apress, 2005) and *SQL Server 2000 Fast Answers for DBAs and Developers* (Apress, 2005). He coauthored *Pro SQL Server 2005* (Apress, 2005) and *Beginning SQL Server 2000 DBA: From Novice to Professional* (Apress, 2004). Joseph graduated with an associate's degree in arts from Bard College at Simon's Rock and earned a bachelor's degree in psychology from the University of Minnesota. You can reach Joseph on his blog, www.joesack.com.

About the Technical Reviewer



EVAN TERRY is the chief technical consultant for The Clegg Company, specializing in data management and information architecture. His past and current clients include the State of Idaho, Albertsons, American Honda Motors, Toyota Motor Sales, The Polk Company, and General Motors. He is the coauthor of Apress's *Beginning Relational Data Modeling*, has published articles in *DM Review*, and has presented at the IAIDQ and DAMA International conferences. For questions or consulting needs, Evan can be contacted at evan_terry@cleggcompany.com.

Acknowledgments

This book is dedicated to David Hatch, and to the family members, friends, and coworkers who helped us get through a very challenging year. From Guillain-Barré syndrome to a broken foot—you were there for us, and we are very lucky to have you in our lives.

During the 9-month writing process, the Apress team helped facilitate a very positive and smooth experience. I want to thank the lead editor, Jonathan Gennick, who was responsive, collaborative, and an all-around great guy to work with. I also appreciate Evan Terry's astute and detailed technical editing—thanks for coming back for a second round!

I also want to thank the amazing Susannah Davidson Pfalzer for her excellent project management skills and positive voice. Thank you also to the keen-eyed Ami Knox, who put the critical finishing touches on this work, and also to Laura Cheu, for the production editing and patience with my last-minute changes.

Lastly—thank you to the rest of the behind-the-scenes Apress team who I may not have met over e-mail or the phone, but who still deserve credit for bringing this book to the market.

Introduction

The purpose of this book is to quickly provide you with the skills you need to solve problems and perform tasks using the Transact-SQL language. I wrote this book in a problem/solution format in order to establish an immediate understanding of a task and its associated Transact-SQL solution. You can use this book to look up the task you want to perform, read how to do it, and then perform the task on your own system. While writing this book, I followed a few key tenets:

- Keep it brief, providing just enough information needed to get the job done.
- Allow recipes and chapters to stand alone—keeping cross-references and distractions to a tolerable minimum.
- Focus on features that are typically implemented entirely using Transact-SQL. For example, I cover the new Resource Governor feature because it will typically be deployed by DBAs using Transact-SQL—whereas I do not cover Policy-Based Management due to its underlying dependencies on SQL Server Agent, SQL Server Management Objects (SMO), and SQL Server Management Studio. Fortunately, most of the new SQL Server engine improvements *are* entirely Transact-SQL based, and therefore are included in this book.
- Write recipes that help a range of skill sets, from novice to professional. I begin each chapter with basic recipes and progressively work up to more advanced topics.

Regarding new SQL Server 2008 features, I have interwoven them throughout the book in the chapters where they apply. If you are just looking for a refresh on new Transact-SQL features, I specifically call them out at the beginning of each chapter in which they exist.

Although a key tenet of this book is to keep things brief, you'll notice that this book is still quite large. This is a consequence of the continually expanding SQL Server feature set; however, rest assured that the recipes contained within are still succinct and constructed in such a way as to quickly give you the answers you need to get the job done.

I've written this book for SQL Server developers, administrators, application developers, and IT generalists who are tasked with developing databases or administering a SQL Server environment. You can read this book from start to finish or jump around to topics that interest you. You can use this book to brush up on topics before a job interview or an exam. Even for the more experienced SQL Server professionals, memory fades—and this book can help quickly refresh your memory on the usage of a command or technique.

Thanks for reading!

